

PART II: Machine Learning: Supervised - Linear Regression

```
In [1]: #Import Python Libraries (NumPy and Pandas)  
import pandas as pd  
import numpy as np
```

```
In [2]: #Import modules and libraries for data visualization  
from pandas.plotting import scatter_matrix  
from matplotlib import pyplot
```

```
In [3]: #Import scikit-Learn module for the algorithm/model (linear regression)  
from sklearn.linear_model import LinearRegression
```

```
In [4]: #Import scikit-Learn module to split the dataset in to train it and test subdatasets  
from sklearn.model_selection import train_test_split
```

```
In [5]: #Import scikit-Learn module for K-fold cross-validation (algorithm/model evaluation and validation)  
from sklearn.model_selection import KFold  
from sklearn.model_selection import cross_val_score
```

```
In [6]: #Import scikit-Learn module classification report to use for information about how the system try to classify  
from sklearn.metrics import classification_report
```

Step 1: Load the data

```
In [ ]:
```

```
In [7]: filename = "/Users/miriamgarcia/Downloads/housing_boston_w_hdrs.csv"
df=pd.read_csv(filename)
print(df)
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	
..	
447	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	
448	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	
449	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	
450	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	
451	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	

	PTRATIO	B	LSTAT	MEDV
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2
..
447	21.0	391.99	9.67	22.4
448	21.0	396.90	9.08	20.6
449	21.0	396.90	5.64	23.9
450	21.0	393.45	6.48	22.0
451	21.0	396.90	7.88	11.9

[452 rows x 14 columns]

Step 2: Preprocess the dataset

```
In [8]: #Clean data and find any missing values
#From looking at the data above we knew that ZN and CHAS had zeros.
#Since most are missing values, it is best to drop them entirely.
df = df.drop("ZN",1)
df = df.drop("CHAS",1)
```

```
In [9]: #Count the number of NaN values in each  
print(df.isnull().sum())
```

```
CRIM      0  
INDUS     0  
NOX       0  
RM        0  
AGE       0  
DIS       0  
RAD       0  
TAX       0  
PTRATIO   0  
B         0  
LSTAT     0  
MEDV      0  
dtype: int64
```

```
In [10]: #Now there is no invalid zero value in any column of the original data.
```

Step 3: Perform the Exploratory Data Analysis (EDA)

```
In [11]: #Get the dimensions/shape of the dataset  
# which will give us the number of records/rows x number of variables/columns  
print(df.shape)
```

```
(452, 12)
```

In [12]: *# Now find the data types of all variables/attributes of the data set*
`print(df.dtypes)`

```
CRIM      float64
INDUS     float64
NOX       float64
RM        float64
AGE       float64
DIS       float64
RAD       int64
TAX       int64
PTRATIO   float64
B         float64
LSTAT     float64
MEDV     float64
dtype: object
```

In [13]: *#Get several records/rows at the top of the dataset, we get 5 to get a feel of the data.*
`print(df.head(5))`

```
      CRIM  INDUS  NOX    RM  AGE   DIS  RAD  TAX  PTRATIO    B  \
0  0.00632  2.31  0.538  6.575  65.2  4.0900    1  296    15.3  396.90
1  0.02731  7.07  0.469  6.421  78.9  4.9671    2  242    17.8  396.90
2  0.02729  7.07  0.469  7.185  61.1  4.9671    2  242    17.8  392.83
3  0.03237  2.18  0.458  6.998  45.8  6.0622    3  222    18.7  394.63
4  0.06905  2.18  0.458  7.147  54.2  6.0622    3  222    18.7  396.90

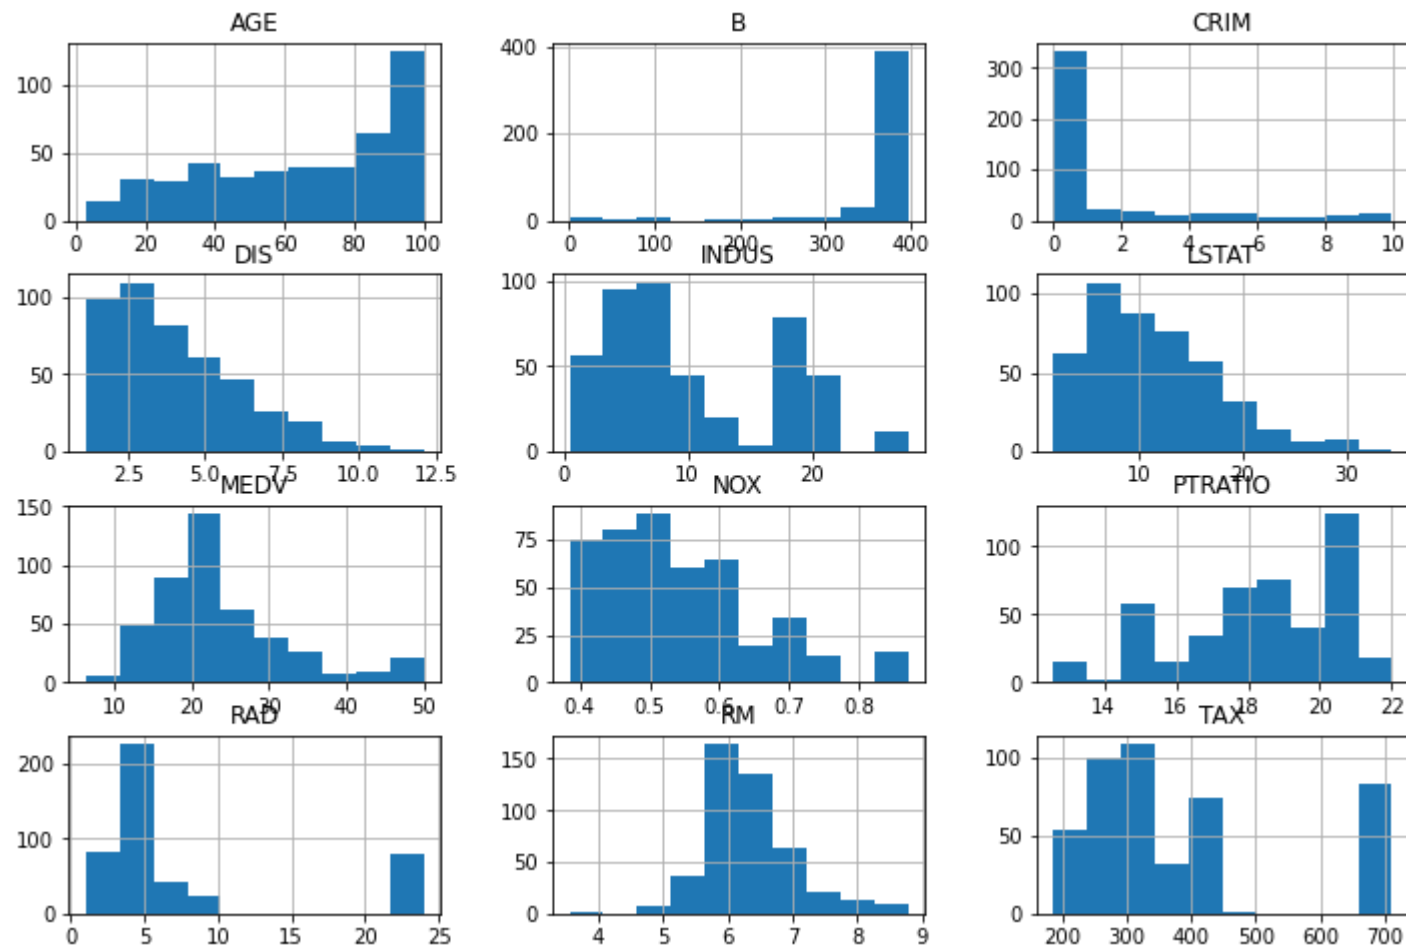
      LSTAT  MEDV
0    4.98  24.0
1    9.14  21.6
2    4.03  34.7
3    2.94  33.4
4    5.33  36.2
```

In [14]: *#We can get the summary statistics of the numeric variables/attributes of the dataset.*
`print(df.describe())`

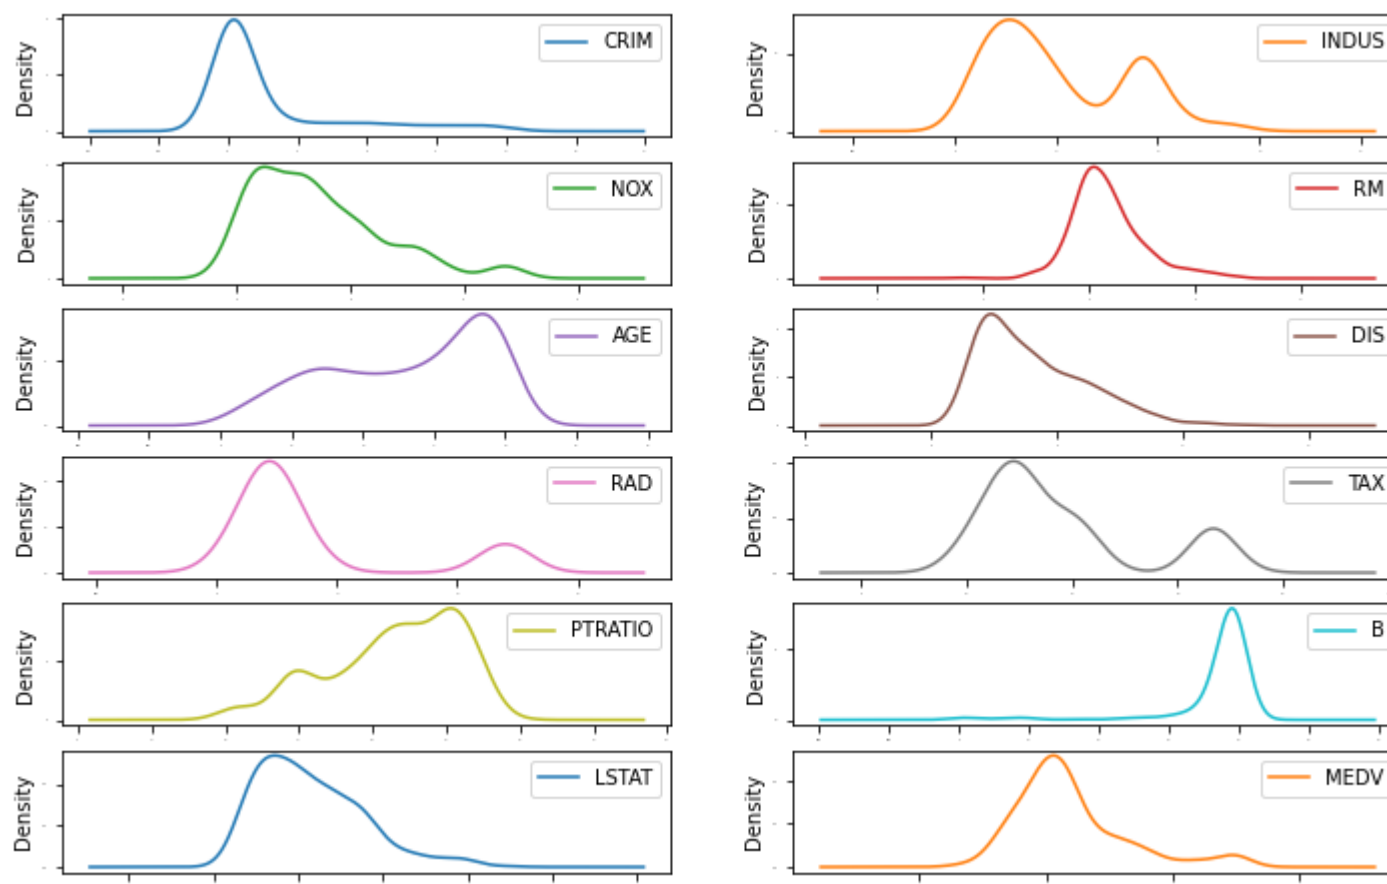
	CRIM	INDUS	NOX	RM	AGE	DIS \
count	452.000000	452.000000	452.000000	452.000000	452.000000	452.000000
mean	1.420825	10.304889	0.540816	6.343538	65.557965	4.043570
std	2.495894	6.797103	0.113816	0.666808	28.127025	2.090492
min	0.006320	0.460000	0.385000	3.561000	2.900000	1.129600
25%	0.069875	4.930000	0.447000	5.926750	40.950000	2.354750
50%	0.191030	8.140000	0.519000	6.229000	71.800000	3.550400
75%	1.211460	18.100000	0.605000	6.635000	91.625000	5.401100
max	9.966540	27.740000	0.871000	8.780000	100.000000	12.126500

	RAD	TAX	PTRATIO	B	LSTAT	MEDV
count	452.000000	452.000000	452.000000	452.000000	452.000000	452.000000
mean	7.823009	377.442478	18.247124	369.826504	11.441881	23.750442
std	7.543494	151.327573	2.200064	68.554439	6.156437	8.808602
min	1.000000	187.000000	12.600000	0.320000	1.730000	6.300000
25%	4.000000	276.750000	16.800000	377.717500	6.587500	18.500000
50%	5.000000	307.000000	18.600000	392.080000	10.250000	21.950000
75%	7.000000	411.000000	20.200000	396.157500	15.105000	26.600000
max	24.000000	711.000000	22.000000	396.900000	34.410000	50.000000

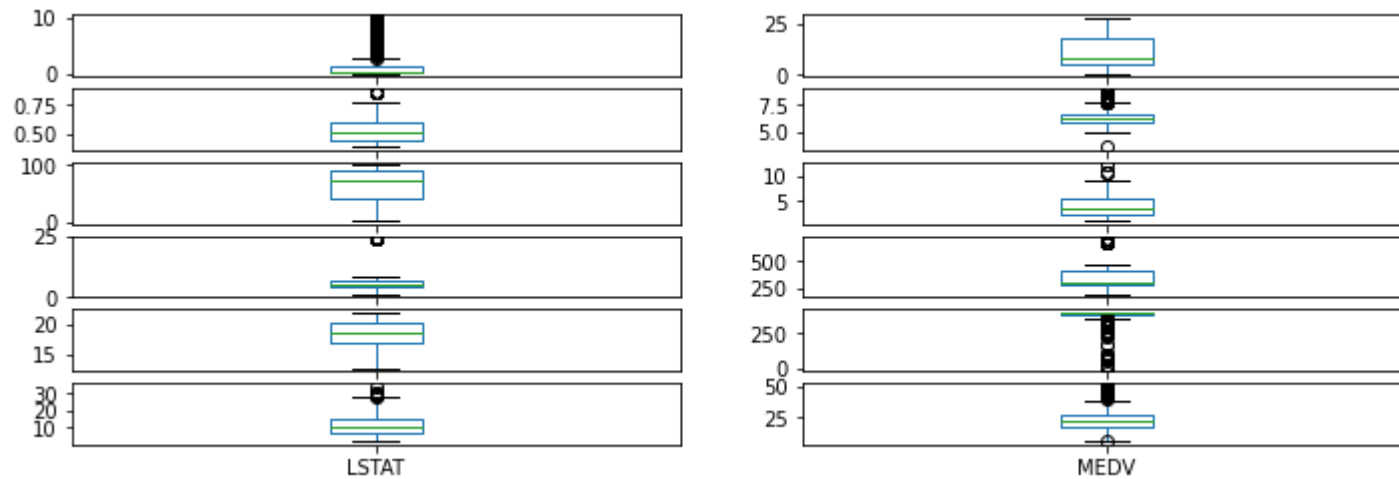
```
In [15]: #Plot histogram for each numeric  
df.hist(figsize=(12, 8))  
pyplot.show()
```



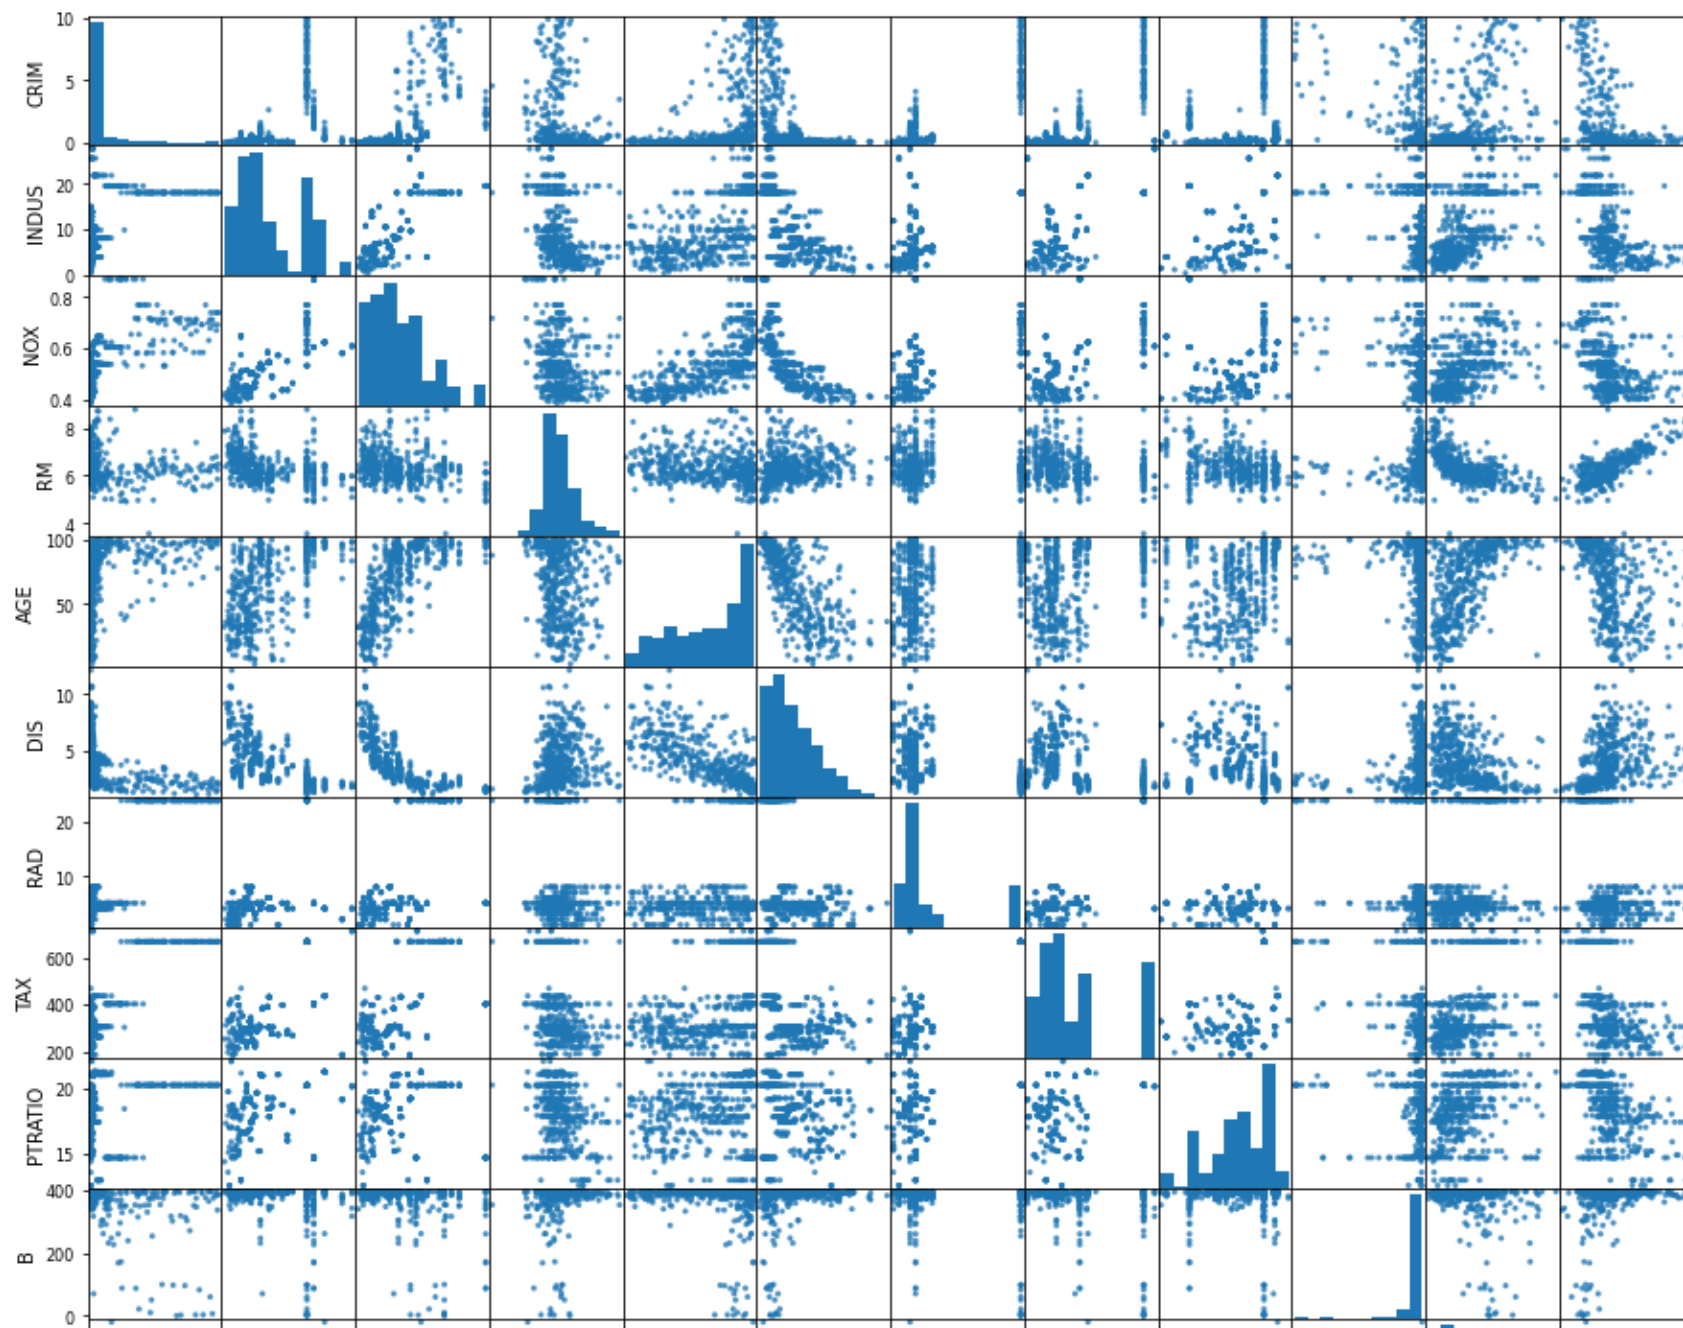
```
In [16]: #Calculate density plots
# 5 numeric variables -> at Least 5 plots -> Layout (2, 3): 2 rows, each row with 3 plots
df.plot(kind='density', subplots=True, layout=(12, 2), sharex=False, legend=True, fontsize=1,
figsize=(12, 16))
pyplot.show()
```




```
In [17]: #Calculate box plots  
df.plot(kind='box', subplots=True, layout=(12,2), sharex=False, figsize=(12,8))  
pyplot.show()
```




```
In [18]: #Calculate scatter plot matrices  
scatter_matrix(df, alpha=0.8, figsize=(15, 15))  
pyplot.show()
```





Step 4: Separate the dataset into the input and output NumPy arrays

```
In [19]: #Then we separate the dataset into input and output NumPy arrays
#We want to store the dataframe values into a NumPy array
array = df.values
#Then we want to separate the array into input and output components by slicing it
#For X (input)[: , 5] --> all the rows, columns from 0 - 4 (5 - 1)
X = array[:,0:11]
#And for Y (output)[: , 5] --> all the rows, column index 5 (Last column)
Y = array[:,11]
```

```
In [20]: print(X)
```

```
[[6.3200e-03 2.3100e+00 5.3800e-01 ... 1.5300e+01 3.9690e+02 4.9800e+00]
 [2.7310e-02 7.0700e+00 4.6900e-01 ... 1.7800e+01 3.9690e+02 9.1400e+00]
 [2.7290e-02 7.0700e+00 4.6900e-01 ... 1.7800e+01 3.9283e+02 4.0300e+00]
 ...
 [6.0760e-02 1.1930e+01 5.7300e-01 ... 2.1000e+01 3.9690e+02 5.6400e+00]
 [1.0959e-01 1.1930e+01 5.7300e-01 ... 2.1000e+01 3.9345e+02 6.4800e+00]
 [4.7410e-02 1.1930e+01 5.7300e-01 ... 2.1000e+01 3.9690e+02 7.8800e+00]]
```

In [21]: `print(Y)`

```
[ 2.31  7.07  7.07  2.18  2.18  2.18  7.87  7.87  7.87  7.87  7.87  7.87
   7.87  8.14  8.14  8.14  8.14  8.14  8.14  8.14  8.14  8.14  8.14  8.14
   8.14  8.14  8.14  8.14  8.14  8.14  8.14  8.14  8.14  8.14  8.14  5.96
   5.96  5.96  5.96  2.95  2.95  6.91  6.91  6.91  6.91  6.91  6.91  6.91
   6.91  6.91  5.64  5.64  5.64  5.64  4.    1.22  0.74  1.32  5.13  5.13
   5.13  5.13  5.13  5.13  1.38  3.37  3.37  6.07  6.07  6.07 10.81 10.81
 10.81 10.81 12.83 12.83 12.83 12.83 12.83 12.83 12.83 4.86 4.86 4.86 4.86
   4.49  4.49  4.49  4.49  3.41  3.41  3.41  3.41 15.04 15.04 15.04 2.89
   2.89  2.89  2.89  2.89  8.56  8.56  8.56  8.56  8.56  8.56  8.56  8.56
   8.56  8.56  8.56 10.01 10.01 10.01 10.01 10.01 10.01 10.01 10.01 10.01
 25.65 25.65 25.65 25.65 25.65 25.65 25.65 21.89 21.89 21.89 21.89 21.89
 21.89 21.89 21.89 21.89 21.89 21.89 21.89 21.89 21.89 21.89 21.89 19.58
 19.58 19.58 19.58 19.58 19.58 19.58 19.58 19.58 19.58 19.58 19.58 19.58
 19.58 19.58 19.58 19.58 19.58 19.58 19.58 19.58 19.58 19.58 19.58 19.58
 19.58 19.58 19.58 19.58 4.05 4.05 4.05 4.05 4.05 4.05 4.05 4.05 2.46
   2.46  2.46  2.46  2.46  2.46  2.46  2.46  3.44  3.44  3.44  3.44  3.44
   3.44  2.93  2.93  0.46  1.52  1.52  1.52  1.47  1.47  2.03  2.03  2.68
   2.68 10.59 10.59 10.59 10.59 10.59 10.59 10.59 10.59 10.59 10.59 10.59
 13.89 13.89 13.89 13.89 6.2   6.2   6.2   6.2   6.2   6.2   6.2   6.2
   6.2   6.2   6.2   6.2   6.2   6.2   6.2   6.2   6.2   6.2   4.93  4.93
   4.93  4.93  4.93  4.93  5.86  5.86  5.86  5.86  5.86  5.86  5.86  5.86
   5.86  5.86  3.64  3.64  3.75  3.97  3.97  3.97  3.97  3.97  3.97  3.97
   3.97  3.97  3.97  3.97  3.97  6.96  6.96  6.96  6.96  6.96  6.96  6.41
   6.41  6.41  6.41  3.33  3.33  3.33  3.33  1.21  2.97  2.25  1.76  5.32
   5.32  5.32  4.95  4.95  4.95 13.92 13.92 13.92 13.92 13.92 13.92 2.24
   2.24  6.09  6.09  6.09  2.18  2.18  2.18  2.18  9.9   9.9   9.9   9.9
   9.9   9.9   9.9   9.9   9.9   9.9   9.9   9.9   7.38  7.38  7.38  7.38
   7.38  7.38  7.38  7.38  3.24  3.24  3.24  6.06  6.06  5.19  5.19  5.19
   5.19  5.19  5.19  5.19  5.19  1.52  1.89  3.78  3.78  4.39  4.39  4.15
   2.01  1.25  1.25  1.69  1.69  2.02  1.91  1.91 18.1  18.1  18.1  18.1
 18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1
 18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1
 18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1
 18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1
 18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1  18.1
 18.1  18.1  27.74 27.74 27.74 27.74 27.74 9.69 9.69 9.69 9.69 9.69
   9.69  9.69  9.69 11.93 11.93 11.93 11.93 11.93]
```

```
In [22]: #Now we want to split the dataset --> training sub-dataset: 67%; and test sub-dataset:
test_size = 0.33
# Selection of records to include in which sub-dataset must be done randomly
#and use this seed for randomization
seed = 10
# Split the dataset (both input & outout) into training/testing datasets
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=test_size,
random_state=seed)
```

Step 5: Split the input/output arrays into the training/testing datasets

```
In [23]: print(X_train)
```

```
[[7.15100e-02 4.49000e+00 4.49000e-01 ... 1.85000e+01 3.95150e+02
 8.44000e+00]
 [8.79212e+00 1.81000e+01 5.84000e-01 ... 2.02000e+01 3.65000e+00
 1.71600e+01]
 [3.51140e-01 7.38000e+00 4.93000e-01 ... 1.96000e+01 3.96900e+02
 7.70000e+00]
 ...
 [6.27390e-01 8.14000e+00 5.38000e-01 ... 2.10000e+01 3.95620e+02
 8.47000e+00]
 [1.69020e-01 2.56500e+01 5.81000e-01 ... 1.91000e+01 3.85020e+02
 1.48100e+01]
 [7.61620e-01 3.97000e+00 6.47000e-01 ... 1.30000e+01 3.92400e+02
 1.04500e+01]]
```

In [24]: `print(Y_train)`

```
[ 4.49 18.1   7.38  3.44 15.04  7.38  8.14  3.97  1.25  8.14 13.92  7.07
   5.13  5.13  8.14  9.9   2.46 13.89 18.1   4.05  8.14  9.69  3.97 19.58
  18.1  10.59 25.65 18.1  18.1   3.33  0.46  9.9   6.41  5.86  5.96 21.89
  18.1   6.41 18.1   3.24 18.1   4.93  2.93 25.65  5.96 18.1   8.14  3.33
   5.19  6.2   3.44  1.52  8.56  2.18  6.06  5.19  5.19  6.2  19.58 18.1
   2.68 18.1   9.69  8.56  6.96  5.86  5.13  4.05  1.52  8.14  5.19  8.56
   4.95 21.89 21.89  6.91 18.1   7.38  5.13 19.58 12.83  7.38 27.74  3.24
  19.58  5.96 13.89  6.2   5.19 18.1  11.93  1.91  5.86 21.89 10.01  9.9
   1.52  6.07  5.86  2.31  5.64  6.41  3.97 18.1   1.22 18.1  18.1   4.15
   4.86  3.97 10.59 18.1  18.1   2.46  4.49  6.91  2.24 19.58  9.69 18.1
  18.1   2.18  3.97 21.89 18.1   1.52  2.46  7.87  3.97  8.14 18.1   3.97
  11.93  8.56  2.95  8.14 10.59 18.1  19.58 18.1  13.92 10.59 19.58  9.69
   8.14  3.64 10.59 21.89 18.1   1.69 19.58  9.69  6.2   6.2  19.58  9.9
   3.97 10.01  2.02  6.91 18.1  18.1  18.1   4.86  2.89  6.07  2.89  8.14
   3.64  5.86 18.1  19.58  4.39 19.58 10.59 18.1 27.74  3.41  6.91  6.2
   5.19 18.1   5.64  4.05  6.96 18.1  18.1   5.64  8.14 18.1   5.19  5.19
   6.96  2.25  2.18  8.14 21.89 19.58  2.03  6.91 18.1   6.2  18.1   6.2
   6.91  6.2  12.83  2.46 18.1   6.2   3.97 10.01  7.87 19.58  6.2   4.49
   5.86 21.89 19.58 18.1   6.96 18.1   4.95  2.46 18.1   2.18 18.1  18.1
  18.1   2.18 21.89  6.91  1.69  5.86  5.32  2.97  2.68 21.89 19.58 18.1
  10.81 10.01 18.1  11.93 18.1  13.89 12.83  9.69 18.1   5.86 15.04 27.74
  15.04  8.14  1.21  4.95  9.9   8.14  7.38 21.89  3.41 19.58  1.76  1.89
   2.01 21.89 18.1  18.1   4.93  8.14 12.83 18.1   2.46  4.05 27.74  1.47
   8.14  5.13 25.65 18.1   3.78  4.   18.1   4.93 18.1   8.14 18.1  18.1
   2.95  3.75 10.81  7.87 11.93 18.1   6.2  19.58 25.65 18.1   7.38  8.14
  25.65  3.97]
```

Step 6: Build and train the model

```
In [25]: #Now we can build the model
model = LinearRegression()
#Then train the model using the training sub-dataset
model.fit(X_train, Y_train)
#Print out the coefficients and the intercept
#Print intercept and coefficients
print (model.intercept_)
print (model.coef_)

-1.1546319456101628e-13
[1.28737288e-14  1.00000000e+00  2.36396189e-14  3.44288400e-15
 1.52655666e-16  3.42087469e-15  5.27355937e-16  1.11022302e-16
 1.37910516e-16  2.77555756e-17  2.22044605e-16]
```

```
In [26]: #We can pair the feature names with the coefficients
#and print out the list with their correspondent variable name
names_2 = ['CRIM', 'INDUS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX', 'PTRATIO', 'B', 'LSTAT', 'MEDV']
```

```
In [27]: coeffs_zip = zip(names_2, model.coef_)
```

```
In [28]: #Convert iterator into set
coeffs = set(coeffs_zip)
```

```
In [29]: #Print (coeffs)
for coef in coeffs:
    print (coef)

('CRIM', 1.287372880437735e-14)
('NOX', 2.3639618908222193e-14)
('DIS', 3.4208746946262636e-15)
('INDUS', 0.9999999999999994)
('B', 2.7755575615628914e-17)
('TAX', 1.1102230246251565e-16)
('PTRATIO', 1.3791051634015616e-16)
('AGE', 1.5265566588595902e-16)
('RM', 3.4428839987277193e-15)
('RAD', 5.273559366969494e-16)
('LSTAT', 2.220446049250313e-16)
```

```
In [30]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

```
Out[30]: LinearRegression(n_jobs=1)
```

Step 7: Calculate the R2 value

```
In [31]: R_squared = model.score(X_test, Y_test)
print(R_squared)
```

```
1.0
```

```
In [32]: #Here we get a perfect R-Squared score of 1. A higher R-Squared
#indicates a higher correlatio between the independent variables and the dependent variables.
```

Step 8: Predict the "Median value of owner-occupied homes in 1000 dollars"

Scenario 1: It is assumed that two new suburbs/towns/developments have been established in the Boston area. The agency has collected the housing data of these two new suburbs/towns/developments.

```
In [33]: #We will use the mean values of the described statistics for the first "made up housing records"
#The suburb area has the following predictors:
#CRIM:1.42
#INDUS:10.30
#NOX:0.54
#RM: 6.34
#AGE:65.56 (proportion of owner-occupied units built prior to 1940)
#DIS: 4.04 (weighted distances to five Boston employment centers)
#RAD: 7.82 (index of accessibility to radial highways)
#TAX: 377.44
#PTRATIO: 18.25 (pupil-teacher ratio by town)
#B:369.83
#LSTAT:23.75
```

```
In [34]: model.predict([[1.42, 10.30, 0.54, 6.34, 65.56,4.04,7.82,377.44,18.25,369.83,23.75]])
```

```
Out[34]: array([10.3])
```

```
In [35]: #The model predicts that the median value of owner-occupied homes
#in 1000 dollars in the above suburb should be around 10,300 under this scenario
```

```
In [36]: #We will now have a "made up housing record" in order to retrain our model.
#The suburb area has the following predictors:
#CRIM:2.8
#INDUS:11.30
#NOX:0.55
#RM: 6.39
#AGE:45.56 (proportion of owner-occupied units built prior to 1940)
#DIS: 8.04 (weighted distances to five Boston employment centers)
#RAD: 4.82 (index of accessibility to radial highways)
#TAX: 277.44
#PTRATIO: 13.25 (pupil-teacher ratio by town)
#B:269.83
#LSTAT:21.75
```

```
In [37]: model.predict([[2.8, 11.30, 0.55, 6.39, 45.56,8.04,4.82,277.44,13.25,269.83,21.75]])
```

```
Out[37]: array([11.3])
```



```
In [38]: #With this second scenario, the model predicts that the median value of owner-occupied homes  
#in 1000 dollars in the above suburb should be around 11,300
```

Step 9: Evaluate the model using the 10-fold cross-validation

```
In [39]: # Evaluate the algorithm  
# Specify the K-size in this case 10-fold  
num_folds = 10  
#We must use the same seed value so that the same subsets can be obtained  
# for each time the process is repeated  
seed = 10  
# Split the whole data set into folds  
kfold = KFold(n_splits=num_folds, random_state=seed)  
# For Linear regression, we can use MSE (mean squared error) value  
# to evaluate the model/algorithm  
scoring = 'neg_mean_squared_error'  
# Train the model and run K-fold cross-validation to validate/evaluate the model  
results = cross_val_score(model, X, Y, cv=kfold, scoring=scoring)  
# Print out the evaluation results  
# Result: the average of all the results obtained from the k-fold cross-validation  
print(results.mean())
```

-1.1773044465875376e-26

/Users/miriamgarcia/opt/anaconda3/lib/python3.8/site-packages/sklearn/model_selection/_split.py:293: FutureWarning: Setting a random_state has no effect since shuffle is False. This will raise an error in 0.24. You should leave random_state to its default (None), or set shuffle=True.

warnings.warn(

```
In [41]: #After we train we evaluate Use K-Fold to determine if the model is acceptable  
#We pass the whole set because the system will divide for us -1.177 average  
#of all errors (mean of square errors)
```

```
In [42]: #
```

```
In [ ]:
```

